Company Overview - Chapter 1

Page 1

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 2

Page 2

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 3

Page 3

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 4

Page 4

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 5

Page 5

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 6

Page 6

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 7

Page 7

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 8

Page 8

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 9

Page 9

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 10

Page 10

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 11

Page 11

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 12

Page 12

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 13

Page 13

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 14

Page 14

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 15

Page 15

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 16

Page 16

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 17

Page 17

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 18

Page 18

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 19

Page 19

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 20

Page 20

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 21

Page 21

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 22

Page 22

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 23

Page 23

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 24

Page 24

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 25

Page 25

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 26

Page 26

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 27

Page 27

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 28

Page 28

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 29

Page 29

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 30

Page 30

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 31

Page 31

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 32

Page 32

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 33

Page 33

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 34

Page 34

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 35

Page 35

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 36

Page 36

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 37

Page 37

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 38

Page 38

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 39

Page 39

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 40

Page 40

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 41

Page 41

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 42

Page 42

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 43

Page 43

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 44

Page 44

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 45

Page 45

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 46

Page 46

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 47

Page 47

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 48

Page 48

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 49

Page 49

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 50

Page 50

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 51

Page 51

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 52

Page 52

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 53

Page 53

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 54

Page 54

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 55

Page 55

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 56

Page 56

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 57

Page 57

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 58

Page 58

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 59

Page 59

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 60

Page 60

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 61

Page 61

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 62

Page 62

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 63

Page 63

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 64

Page 64

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 65

Page 65

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 66

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 67

Page 67

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 68

Page 68

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 69

Page 69

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 70

Page 70

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 71

Page 71

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 72

Page 72

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 73

Page 73

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 74

Page 74

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 75

Page 75

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 76

Page 76

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 77

Page 77

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 78

Page 78

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 79

Page 79

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 80

Page 80

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 81

Page 81

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 82

Page 82

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 83

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 84

Page 84

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 85

Page 85

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 86

Page 86

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 87

Page 87

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 88

Page 88

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 89

Page 89

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 90

Page 90

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Company Overview - Chapter 91

Page 91

The company develops enterprise software, cloud solutions, AI applications, cybersecurity services, and business analytics platforms.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

HR Policies - Chapter 92

Page 92

Attendance, leave, remote work, code of conduct, performance reviews, promotions, payroll, compensation, travel policy, onboarding, offboarding.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

IT Policies - Chapter 93

Page 93

Password policy, MFA, VPN usage, endpoint security, acceptable use, email policy, data classification, backup policy.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Software Development - Chapter 94

Page 94

Agile, Scrum, Git workflow, pull requests, branching strategy, CI/CD, testing, deployment, documentation.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cybersecurity - Chapter 95

Page 95

OWASP Top 10, phishing awareness, incident response, secure coding, vulnerability management.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Cloud - Chapter 96

Page 96

AWS fundamentals, IAM, EC2, S3, Lambda, networking, monitoring, billing.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

AI & Data - Chapter 97

Page 97

Machine Learning lifecycle, RAG, vector databases, embeddings, LLMs, evaluation, prompt engineering.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Business Analysis - Chapter 98

Page 98

Requirements gathering, UML, BPMN, stakeholder analysis, user stories, acceptance criteria.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

Project Management - Chapter 99

Page 99

Risk management, project planning, estimation, communication, reporting.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions
1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:
This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.

FAQs - Chapter 100

Page 100

Frequently asked questions about policies, technical standards, and employee benefits.

This page is intentionally created as part of a realistic enterprise knowledge base for Retrieval-Augmented Generation (RAG) training. Students will use this handbook to practice: Document loading Text extraction Chunking strategies (200, 500, 1000 tokens) Chunk overlap experiments Embedding generation Vector indexing Semantic search Metadata filtering Hybrid search Reranking LLM prompting RAG evaluation using Recall, Precision, Hit@K, MRR, Faithfulness, and Answer Correctness. Scenario Employees frequently ask questions about company policies. The assistant must retrieve the relevant section before generating an answer. Similar concepts may appear across multiple chapters, requiring the retriever to locate the best context.

Example Questions

1. What is the leave policy?
2. How is performance evaluated?
3. What are the password requirements?
4. Explain Git branching strategy.
5. What is the difference between RAG and fine-tuning?
6. When should MFA be enabled?
7. How are internships selected?
8. What is the promotion process?
9. Describe the CI/CD pipeline.
10. What are the responsibilities of a Business Analyst?

Additional Notes:

This content intentionally repeats related terminology with variations to create realistic semantic retrieval challenges. Students should compare embedding models, vector databases, chunk sizes, overlap values, Top-K retrieval, and prompt templates while measuring retrieval quality.